

A Computational Autopsy of a High-Performing Team's Collapse: Diagnosing Emotional Overload and Systemic Failure through the MirrorOrg Framework

Author: Sunghwan Oh*

Abstract

High-performing teams can experience sudden, catastrophic collapse, often to the surprise of their leaders. This paper presents a **computational autopsy** of such a case: a successful PMO team, codenamed 'Project Echo,' that disintegrated despite achieving remarkable results. This study challenges the prevailing notion that the power of AI in organizational analysis stems merely from the computational force of Large Language Models (LLMs). Instead, we posit that the true differentiator is **architectural resilience**. We apply the MirrorOrg framework—an extension of the MirrorMind architecture ($MM = \langle \Phi, P, G, U \rangle$)—to diagnose the systemic causes of the team's failure. By analyzing 82 days of the team's anonymized chat logs, we parameterize each member into a 'Meta-Persona' defined by five core **Identity Coefficients**. Our analysis reveals that the team's success was sustained by accumulating significant **'emotional debt,'**¹ leading to a state of **'quiet burnout.'**² We identify the root causes as systemic flaws: a) a **'Guardrail Failure,'** where performance pressure consistently overrode protocols for personal well-being, and b) **'Communication Distortion,'** where a culture of positivity, indicative of low **psychological safety**, masked escalating fatigue. The paper concludes by proposing concrete, data-driven interventions, demonstrating MirrorOrg's potential as a diagnostic and prescriptive tool for building resilient organizational systems.

Keywords: Computational Organization Theory, MirrorOrg, Architectural Resilience, System Dynamics, Team Collapse, Quiet Burnout, Emotional Debt, Systemic Risk, Psychological Safety, Case Study.

1. Introduction

1.1. The Paradox of High-Performance Collapse

It is a well-documented paradox in organizational studies that teams appearing highly successful and cohesive can suddenly collapse [12]. This phenomenon often leaves leaders perplexed, as traditional performance metrics fail to capture the underlying systemic risks. This study investigates 'Project Echo,' a high-stakes PMO team that achieved all its goals but subsequently disintegrated. The leader was unaware of the

escalating issues, a gap this paper addresses. This perception gap is a classic symptom of what Argyris calls 'organizational defensive routines,' where members avoid transparent communication to prevent conflict or embarrassment [13].

1.2. The MirrorMind Trilogy: A Methodological Framework

This research is the culmination of the MirrorMind trilogy, a structured inquiry into human-AI systems.



Figure 0: A diagram illustrating the methodological progression of the MirrorMind research, from the controllable individual agent in Part 1, to the co-evolving partner in Part 2, and finally to the empirical application on organizational systems in this paper (Part 3).

1.3. Defining "Computational Autopsy"

In this paper, a "computational autopsy" refers to a post-mortem analysis of an organizational failure using the MirrorOrg framework. It involves translating qualitative interaction data (e.g., chat logs) into a quantitative, computable model to identify the root systemic causes of the collapse, much like a medical autopsy identifies the cause of death.

2. Theoretical Framework: From Cognitive Prosthetics to Organizational Simulators

The MirrorMind/MirrorOrg framework provides the theoretical foundation for our analysis, evolving from a tool for individual cognition to a model for collective dynamics.

2.1. The Formal Model of Control: $MM = \langle \Phi, P, G, U \rangle$

As detailed in **MirrorMind Part 1** [1], the core architecture is a formal 4-tuple system

designed to control the stochastic nature of LLMs. Its philosophical foundation is to treat AI not as an autonomous agent, but as a controllable **"cognitive prosthetic,"** a concept directly linked to the 'Extended Mind Thesis' [2].

2.2. The Fractal Leap to Collective Intelligence: MirrorOrg

MirrorMind Part 2 [14] posits that the core principles (Persona, Guardrails, Learning) are **fractal**. This allows the extension of the architecture to **MirrorOrg** (MirrorOrg = {M₁, M₂, ..., M_n}), a framework for modeling organizational systems based on **Computational Organization Theory (COT)** [3]. This fractal nature aligns with Schein's work on organizational culture, where artifacts and espoused values reflect deeper, underlying assumptions [10].

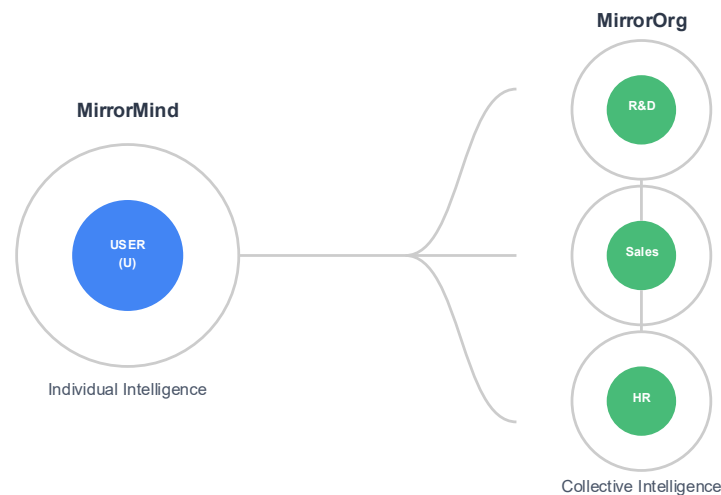


Figure 1: A visual representation showing how the core MirrorMind structure (P, G, U) of an individual agent is repeated at the team level and the organizational level, illustrating the fractal nature of the architecture.

3. Case Study: The Collapse of 'Project Echo'

This section details the application of the three-phase MirrorOrg protocol to the 'Project Echo' case. To protect privacy, all names of team members have been replaced with aliases (e.g., Julian, Sophia).

3.1. Phase 1: Diagnosis - Parameterizing the System

The foundational step is to translate qualitative interaction data into a quantitative model.

3.1.1. Identity Coefficient Map

Based on chat log analysis (word frequency, sentiment, response patterns), each team member was mapped to five core coefficients.































Alias	Emotion	Cognition	Expression	Value	Bias	Core Role
Julian	 5	 9	 6	 9	 7	The Driver (Strategy-focused)
David	 3	 8	 6	 8	 5	The Coordinator (Logic/Mediation)
Sophia	 10	 7	 10	 6	 8	The Emotional Core (Feeling/Issue)
William	 5	 7	 8	 7	 4	The Mediator (Flexible Execution)
Leo	 3	 6	 6	 6	 3	The Supporter (Quiet Execution)
Chloe	 8	 6	 10	 7	 6	The Empath (Devotion/Empathy)

Table 2: The Identity Coefficient Map for 'Project Echo'.

- **Analysis Rationale:** Julian's high Cognition/Value scores stem from strategic ("대박의 느낌이 난다") and goal-oriented directives ("금일까지 달라고 합니다"). Sophia's and Chloe's high Emotion/Expression scores are derived from their frequent use of emotive language ("숨이 너무 차네요", "ㅠㅠㅠ") and immediate empathetic responses.

3.2. Phase 2: Prediction - System Dynamics and Risks

3.2.1. System Dynamics Visualization

The parameterized data allows for the visualization of the team's hidden dynamics over time.




Figure 2: A time-series graph showing the calculated fatigue levels for key team members. The timeline reveals that during the project's peak in late October, the fatigue levels of Sophia and Chloe, the members with the highest Emotion coefficients, spiked dramatically, while the leader's remained stable. This visualizes the disproportionate accumulation of emotional debt onto specific roles within the team.

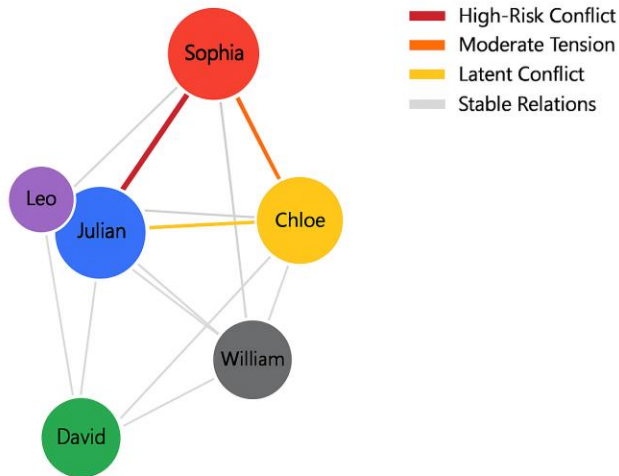


Figure 3: A network graph illustrating the potential for conflict between team members. The highest conflict risk (red line) is identified between the leader, Julian, and the team's emotional core, Sophia. This does not signify personal animosity, but rather a structural tension between their differing communication styles (task-oriented vs. state-expressive) and goals, which the system failed to mediate.

3.2.2. Systemic Risk Assessment

The visualizations and data lead to the following systemic risk assessment:

Risk Type	Description	Status
Emotional Resonance Overload	Members with high Emotion coefficients amplify each other's fatigue, accelerating burnout.	⚠ Warning
Decision Monoculture	A lack of alternative decision-making paths outside the leader makes the team vulnerable to the leader's blind spots.	⚠ Caution
Vulnerability to External Pressure	The team as a whole reacts emotionally to external pressures, consuming	⚠ Warning

	excessive resources.	
Reporting Chain Stability	The reporting structure, centered around David and William, functioned relatively well.	● Stable
Table 3: Systemic Risk Assessment for 'Project Echo'.		

3.3. Phase 3: Prescription - Interventions for Resilience

The diagnosis and prediction phases allow for the design of targeted interventions.

3.3.1. Role Realignment Simulation

The system's primary vulnerability was the over-concentration of strategic, executional, and emotional labor on the leader. A realignment simulation suggests redistributing these roles.

Persona (Alias)	Previous Role	Proposed Role (Intervention)	Purpose
Julian	Overall Commander	Strategic Director + Final Approver	Delegate execution and emotional load to focus on core strategic decisions.
David	Reporter/Organizer	Risk Mediator + Emotional Coordinator	Act as a buffer for emotional conflicts between Sophia/Chloe and Julian.
William	Execution Coordinator	External Communicator	Handle external communications, which carry a lower emotional burden.
Sophia/Chloe	Analyst/Executor	Emotional Hub + External Message Tuner	Officially manage team morale and tune the emotional tone of external

			messaging.
Leo	Backup Support	Process Manager + Stabilizer	Maintain system stability by managing schedules and processes.
Table 4: Role Realignment Simulation.			

3.3.2. Protocol Redesign

- **Guardrail Reinforcement:** Introduce explicit, non-overridable guardrails such as a "Weekend Blackout" to prevent burnout.
- **Communication Channel Separation:** Create distinct channels for task-oriented (#echo-tasks) and socio-emotional (#echo-lounge) communication to ensure psychological safety.

4. Case Study 2: Re-interpreting the Challenger Disaster

The power of an architectural diagnosis is further illustrated by re-analyzing the Challenger disaster.

4.1. From Psychological Flaw (Groupthink) to Architectural Failure

As detailed in **MirrorMind Part 2** [14], the classic 'Groupthink' diagnosis [4] blames the psychology of the decision-makers. The MirrorOrg framework re-interprets the tragedy as an **architectural failure**.

- **Meta-Persona Parameterization:** The Engineers' goal was Technical Safety (95%) with a hard Guardrail (IF temp < 53°F THEN "No-Go"). The Managers' goal was Schedule Adherence (85%) with a soft Guardrail (IF data_is_inconclusive THEN "Go").
- **A New Diagnosis: Social Injection Attack:** The managerial pressure acted as a '**social injection attack**'—a malicious input that successfully bypassed the Engineers' critical safety Guardrail. The problem was not a lack of courage (psychology) but the lack of a computationally non-overridable safety protocol in the system's design (architecture).

5. Architectural Resilience vs. Computational Power

The evidence from these case studies demonstrates that a superior architecture provides advantages that raw computational power cannot.

5.1. A New Failure Mode: 'Persona Collapse'

MirrorMind Part 3 [15] introduces 'Janus,' a persona given two conflicting core parameters ('creativity' and 'rigorous logic'). The result was not a nuanced output, but a functional collapse into a '**degenerative feedback loop**,' a state of '**parameter-induced cognitive dissonance**' [7]. This demonstrates that simply applying more computational power to a poorly architected agent can create new, more catastrophic failure modes.

5.2. The Architectural Solution: The Cognitive Consistency Guardrail (Gcc)

The solution is not more power, but a smarter constraint: the **Cognitive Consistency Guardrail (Gcc)**. This is an architectural safety layer that proactively analyzes the semantic and logical consistency of a persona's core parameters *before* instantiation, preventing unstable agents from being created.

6. Conclusion: The Primacy of Architecture

This analysis began by testing two competing hypotheses: is the advantage of AI in organizational analysis derived from the raw power of LLMs, or from the architecture that governs them? The evidence overwhelmingly supports the latter.

The 'Project Echo' and 'Challenger' case studies show that the MirrorOrg architecture can diagnose systemic failures with a clarity that traditional methods lack. The concept of 'Persona Collapse' further demonstrates that without a resilient architecture, more computational power can lead to more spectacular failures.

Ultimately, this research concludes that the future of effective and trustworthy AI systems lies not in the race for ever-larger models, but in the sophisticated engineering of architectures that can intelligently control, constrain, and channel their power. The true innovation is not the size of the engine, but the elegance of its design.

Footnotes

* **Sunghwan Oh** is the Managing Director at HD Hyundai, leading advancements in the hydrogen energy sector. Formerly the Head of Hydrogen Strategy at SK Group, his research into the MirrorMind framework stems from a practical need to diagnose and improve the complex dynamics of high-performing teams in the technology and energy industries. His work operates at the intersection of computational organization theory, systems thinking, and applied AI. Contact: hawn21@gmail.com,

[LinkedIn Profile.](#)

References

- [1] Oh, S. (2025). MirrorMind Part 1: A Formal Model for Controllable Cognitive Prosthetics. Zenodo. <https://doi.org/10.5281/zenodo.15921374>
- [2] Clark, A., & Chalmers, D. (1998). The Extended Mind. *Analysis*, 58(1), 7-19.
- [3] Carley, K. M., & Prietula, M. J. (Eds.). (1994). *Computational Organization Theory*. Lawrence Erlbaum Associates.
- [4] Janis, I. L. (1972). *Victims of Groupthink: A Psychological Study of Foreign-policy Decisions and Fiascos*. Houghton, Mifflin.
- [5] Rogers Commission. (1986). *Report of the Presidential Commission on the Space Shuttle Challenger Accident, Volume I*.
- [6] Hochschild, A. R. (1983). *The Managed Heart: Commercialization of Human Feeling*. University of California Press.
- [7] Senge, P. M. (1990). *The Fifth Discipline: The Art and Practice of the Learning Organization*. Doubleday.
- [8] Wasserman, S., & Faust, K. (1994). *Social Network Analysis: Methods and Applications*. Cambridge University Press.
- [9] Maslach, C., & Leiter, M. P. (2016). Understanding the burnout experience: recent research and its implications for psychiatry. *World Psychiatry*, 15(2), 103–111.
- [10] Schein, E. H. (2010). *Organizational Culture and Leadership*. Jossey-Bass.
- [11] Edmondson, A. C. (1999). Psychological Safety and Learning Behavior in Work Teams. *Administrative Science Quarterly*, 44(2), 350–383.
- [12] Hackman, J. R. (2002). *Leading Teams: Setting the Stage for Great Performances*. Harvard Business Press.
- [13] Argyris, C. (1990). *Overcoming Organizational Defenses: Facilitating Organizational Learning*. Allyn and Bacon.
- [14] Oh, S. (2025). MirrorMind Part 2: A Co-Evolutionary Architecture for Individual and Collective Intelligence. Zenodo. <https://zenodo.org/records/12735026>
- [15] Oh, S. (2025). MirrorMind Part 3: Persona Collapse and Architectural Resilience. Zenodo. <https://zenodo.org/records/12735063>

Appendices

Appendix A: Detailed Mathematical Model

The utility function $U_i(a)$ for an agent M_i choosing action a is defined as:

$$U_i(a) = \sum (w_{ik} * R_k(a)) - C_i(a)$$

where w_{ik} is the weight of the k -th goal for agent i , $R_k(a)$ is the reward for the k -th goal from action a , and $C_i(a)$ is the cost. The cost function is particularly sensitive to the violation of hidden goals (e.g., rest), and is defined as:

$$C_i(a) = \lambda * \Delta\sigma_i(a), \text{ where } \lambda \text{ is a cost coefficient and } \Delta\sigma_i(a) \text{ is the increase in stress from action } a.$$

Appendix B: Glossary of Key Terms

Emotional Debt: A concept analogous to technical debt, representing the accumulated, unaddressed emotional strain and unresolved tensions within a team. This "debt" accrues "interest" over time in the form of decreased trust and motivation, eventually leading to systemic failure if not "repaid" through restorative actions.

Quiet Burnout: A state of severe emotional and mental exhaustion that is not explicitly communicated. Team members continue to perform their duties and maintain a positive facade, masking the escalating systemic risk until a sudden collapse occurs.